



Università
di Genova

DIBRIS DIPARTIMENTO
DI INFORMATICA, BIOINGEGNERIA,
ROBOTICA E INGEGNERIA DEI SISTEMI

Weather Prediction

Machine Learning and Data Analysis

Samuele Crea (4868097) & Marco Zucca (4828628)

Data Analysis

Exploration, cleaning and analysis of the dataset

Dataset exploration and cleaning

Dataset exploration and cleaning

Slice of the original dataset

| | DATE | MONTH | BASEL_cloud_cover | BASEL_humidity | BASEL_pressure | BASEL_global_radiation | BASEL_precipitation | BASEL_sunshine | BASEL_temp_mean | BASEL_temp_min | ... | STOCKHOLM_temp_min |
|---|----------|-------|-------------------|----------------|----------------|------------------------|---------------------|----------------|-----------------|----------------|-----|--------------------|
| 0 | 20000101 | 1 | 8 | 0.89 | 1.0286 | 0.20 | 0.03 | 0.0 | 2.9 | 1.6 | ... | -9.3 |
| 1 | 20000102 | 1 | 8 | 0.87 | 1.0318 | 0.25 | 0.00 | 0.0 | 3.6 | 2.7 | ... | 0.5 |
| 2 | 20000103 | 1 | 5 | 0.81 | 1.0314 | 0.50 | 0.00 | 3.7 | 2.2 | 0.1 | ... | -1.0 |
| 3 | 20000104 | 1 | 7 | 0.79 | 1.0262 | 0.63 | 0.35 | 6.9 | 3.9 | 0.5 | ... | 2.5 |
| 4 | 20000105 | 1 | 5 | 0.90 | 1.0246 | 0.51 | 0.07 | 3.7 | 6.0 | 3.8 | ... | -1.8 |

- Includes measurements of meteorological data from 18 different European cities
- obtained for each day from 2000 to 2009 included
- Several collected features that allow the study of weather

Dataset exploration and cleaning

Features per city

| Feature (type) | Column name | Description | Physical Unit |
|------------------|------------------|------------------------|---------------|
| mean temperature | temp_mean | mean daily temperature | in 1 °C |
| max temperature | temp_max | max daily temperature | in 1 °C |
| min temperature | temp_min | min daily temperature | in 1 °C |
| cloud_cover | cloud_cover | cloud cover | oktas |
| global_radiation | global_radiation | global radiation | in 100 W/m2 |
| humidity | humidity | humidity | in 1 % |
| pressure | pressure | pressure | in 1000 hPa |
| precipitation | precipitation | daily precipitation | in 10 mm |
| sunshine | sunshine | sunshine hours | in 0.1 hours |
| wind_gust | wind_gust | wind gust | in 1 m/s |
| wind_speed | wind_speed | wind speed | in 1 m/s |

Dataset exploration and cleaning

Cleaning

```
Missing values per column:
```

```
DATE          0
MONTH         0
BASEL_cloud_cover  0
BASEL_humidity  0
BASEL_pressure  0
```

```
..
```

```
TOURS_global_radiation  0
TOURS_precipitation     0
TOURS_temp_mean         0
TOURS_temp_min          0
TOURS_temp_max          0
```

```
Length: 165, dtype: int64
```

```
Number of duplicate rows: 0
```

There are no missing values or duplicate rows

We don't need any additional dataset filling actions

Dataset exploration and cleaning

Cleaning

- **We decided to take in account only one city: Munich**
 - By dropping every column not related to that city
 - Renaming the features to not include the city name as prefix
- **Converting the 'DATE' feature**
 - from YYYYMMDD to three different features: 'year', 'month' and 'day'
 - the 'DATE' feature dropped
- **Creating new boolean features (features engineering)**
 - 'rainy_day' based on 'precipitation'
 - 'sunny_day' based on 'sunshine'
 - 'good_day' based on 'precipitation', 'sunshine' and 'temp_mean'

Dataset exploration and cleaning

Cleaned dataset

| | day | month | year | cloud_cover | wind_speed | wind_gust | humidity | pressure | global_radiation | precipitation | sunshine | temp_mean | temp_min | temp_max | rainy_day | sunny_day | good_day |
|---|-----|-------|------|-------------|------------|-----------|----------|----------|------------------|---------------|----------|-----------|----------|----------|-----------|-----------|----------|
| 0 | 1 | 1 | 2000 | 8 | 2.6 | 9.4 | 0.91 | 1.0273 | 0.20 | 0.20 | 0.0 | 1.7 | -0.5 | 2.6 | 1 | 0 | 0 |
| 1 | 2 | 1 | 2000 | 6 | 2.1 | 8.2 | 0.90 | 1.0321 | 0.66 | 0.00 | 6.1 | 1.9 | -0.2 | 5.8 | 0 | 1 | 0 |
| 2 | 3 | 1 | 2000 | 7 | 2.1 | 6.9 | 0.92 | 1.0317 | 0.28 | 0.00 | 0.4 | -0.4 | -3.3 | 0.9 | 0 | 0 | 0 |
| 3 | 4 | 1 | 2000 | 6 | 2.7 | 11.7 | 0.75 | 1.0260 | 0.58 | 0.04 | 4.5 | 3.8 | -2.8 | 6.6 | 1 | 0 | 0 |
| 4 | 5 | 1 | 2000 | 5 | 3.3 | 13.2 | 0.87 | 1.0248 | 0.26 | 0.00 | 0.2 | 5.3 | 4.3 | 7.3 | 0 | 0 | 0 |

This will be the dataset used for our studies

Aim of the project

Aim of the project

Our study

Our study focuses on predicting whether or not it will rain based on the features available to us




Data Visualization

Data Visualization

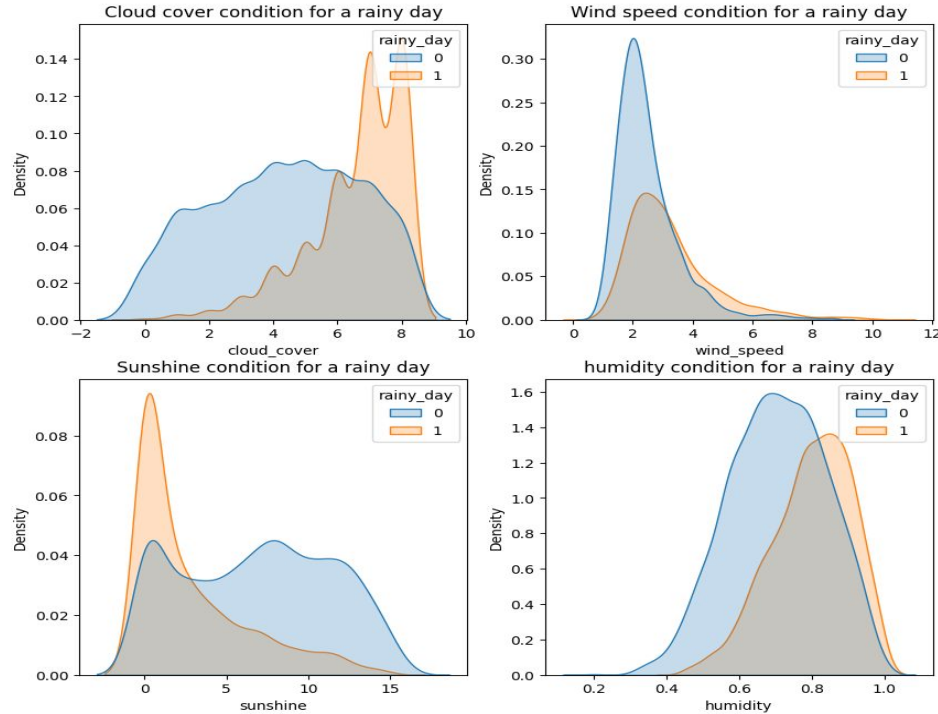
Correlation between the features

- calculated using Pearson coefficient
- Some meaningful correlations:
 - temp_mean/global_radiation
 - cloud_cover/humidity
 - cloud_cover/rainy_day





Data Visualization

Rainy day based on other features



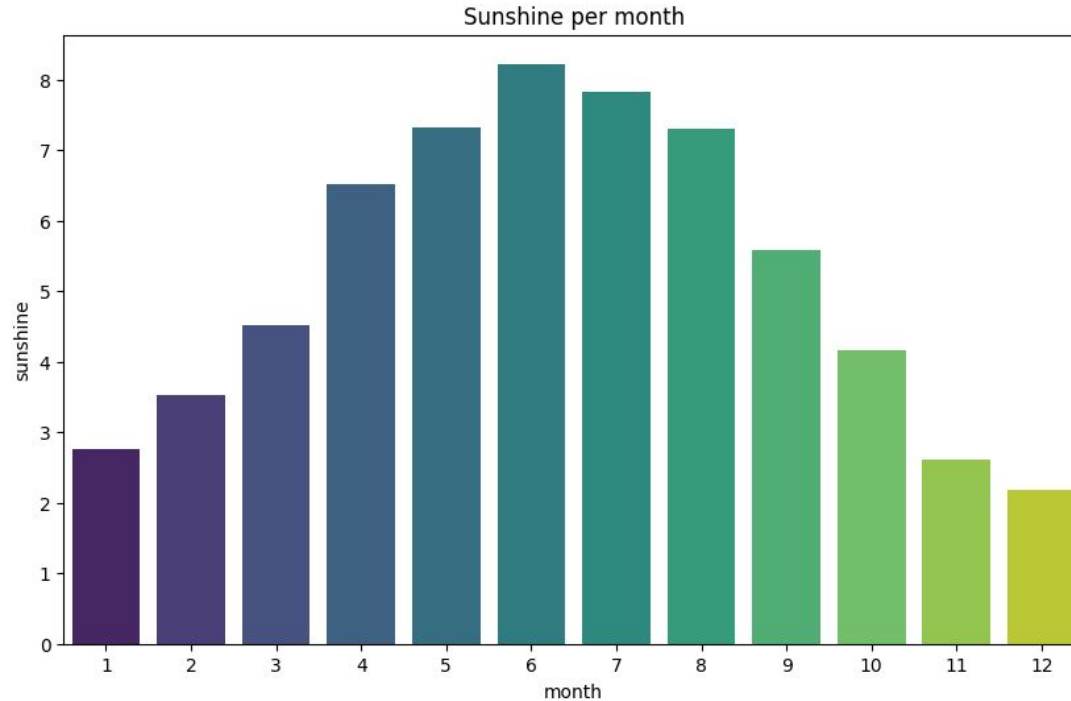
Data Visualization

precipitation and temperature curves



Data Visualization

The sunshine during the months



Machine Learning

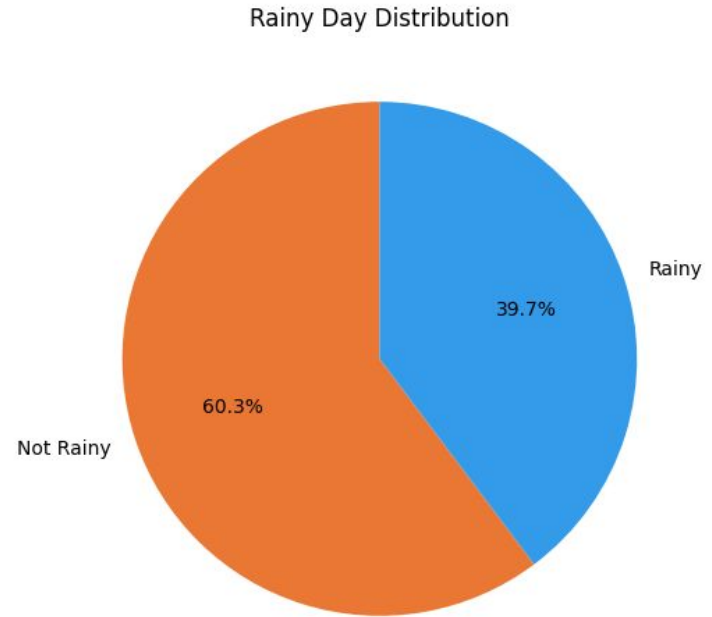
Study on different machine learning methods

Preliminary actions for machine learning

Preliminary actions for machine learning

Data balancing

- the data taken for the training of the dataset are balanced
- not necessary to undersample the data
- Data scaled between the interval $[0,1]$ with `MinMaxScaler()`



Preliminary actions for machine learning

KFold

- Used for the creation of train and test sets for machine learning
 - x sets with every features except rainy_day and precipitation
 - y sets with only boolean values (0,1) from the rainy_day feature
- Is a cross-validator function that provides train/test indices to split data in train/test sets, Split dataset into k consecutive folds
- Improve hyper-parameters tuning

Preliminary actions for machine learning

The metrics

| | |
|-----------------|--|
| Accuracy | Measures the proportion of correct predictions to the total number of predictions |
| Roc-Auc | Used to evaluate the performance of binary classification models. A higher ROC AUC score indicates that the model is better at distinguishing between positive and negative cases. |
| f1 | It is the number of true positive predictions divided by the sum of the true positive predictions and false negative predictions. Recall is an important metric when it is important to minimize the number of false negatives. |
| recall | The harmonic mean of precision and recall, used when there is a need to balance the trade-off between precision and recall. |

The models


The models

SVC

A support vector machine (SVM) is a supervised machine learning model that uses classification algorithms for two-group classification problems.

Results of test set prediction analyzing:

Accuracy of prediction: 83.56%
roc-auc score: 83.59%
f1 score: 83.15%
recall score: 84.33%



The models

Tuned SVC


Results of test set prediction analyzing:

Accuracy of prediction: 83.29%

roc-auc score: 83.33%

f1 score: 82.91%

recall score: 84.33%



The models

SVC TUNING

| Hyperparameters | Brief description | Values to be selected | Default values | Selected values |
|-----------------|---|----------------------------|----------------|-----------------|
| C | Regularization parameter | [0.1,0.5,1,5] | 1.0 | 0.5 |
| gamma | Kernel coefficient | ['scale','auto'] | 'scale' | 'scale' |
| kernel | the kernel type to be used in the algorithm | ['rbf', 'poly', 'sigmoid'] | 'rbf' | 'rbf' |

The models


Logistic Regression

Logistic regression is a statistical analysis method to predict a binary outcome based on prior observations of a data set.

It predicts a dependent data variable by analyzing the relationship between one or more existing independent variables.

Results of test set prediction analyzing:

Accuracy of prediction: 81.37%
roc-auc score: 81.41%
f1 score: 80.95%
recall score: 82.34%



The models

Tuned Logistic Regression


Results of test set prediction
analyzing:

Accuracy of prediction: 81.51%

roc-auc score: 81.54%

f1 score: 81.07%

recall score: 82.34%



The models

Logistic Regression

| Hyperparameters | Brief description | Values to be selected | Default Values | Selected values |
|-----------------|--|---|----------------|-----------------|
| C | Regularization intensity | <code>np.logspace(-3, 3, 10)</code> | 1.0 | 0.1 |
| penalty | Penalty type | <code>["l1", "l2", "elasticnet", None]</code> | 'l2' | l2 |
| solver | Algorithm to use in the optimization problem | <code>["lbfgs", "saga", "sag"]</code> | 'lbfgs' | saga |

The models


KNN

KNN is an algorithm that can be used to solve both classification and regression problems.

it's used in pattern recognition for the classification of objects based on the characteristics of the objects close to the one considered.

Results of test set prediction analyzing:

Accuracy of prediction: 78.9%
roc-auc score: 79.01%
f1 score: 78.85%
recall score: 81.77%



The models

Tuned KNN


Results of test set prediction
analyzing:

Accuracy of prediction: 83.29%

roc-auc score: 83.43%

f1 score: 83.38%

recall score: 87.18%



The models

KNN TUNING

| Hyperparameters | Brief description | Values to be selected | Default Value | Selected values |
|-----------------|------------------------------------|---------------------------------------|---------------|-----------------|
| metric | distance computation metrics | ['minkowski','euclidean','chebyshev'] | 'minkowski' | euclidean |
| n_neighbors | Number of neighbors | np.arange(2,25) | 5 | 24 |
| weights | Weight function used in prediction | ['uniform', 'distance'] | 'uniform' | distance |


The models

Random Forest

Random forest is a learning method for classification and regression that operates by constructing a multitude of decision trees at training time, a decision tree is a graph of decisions

Results of test set prediction analyzing:

Accuracy of prediction: 82.88%
roc-auc score: 82.91%
f1 score: 82.47%
recall score: 83.76%




The models

Tuned Random Forest

Results of test set prediction
analyzing:

Accuracy of prediction: 82.74%
roc-auc score: 82.82%
f1 score: 82.55%
recall score: 84.9%



The models

Random Forest Tuning

| Hyperparameters | Brief description | Values to be selected | Default Value | Selected values |
|-------------------|--|------------------------|---------------|-----------------|
| max_features | number of features to consider when looking for the best split | ['sqrt', 'log2', None] | 'sqrt' | log2 |
| max_depth | The maximum depth of the tree | [10, 30, 60, None] | None | 60 |
| min_samples_leaf | The minimum number of samples required to be at a leaf node. | [1, 2, 4] | 1 | 4 |
| min_samples_split | The minimum number of samples required to split an internal node | [2,5] | 2 | 5 |
| n_estimators | number of trees in the forest | [50,100, 200, 300] | 100 | 100 |
| bootstrap | Whether bootstrap samples are used when building trees | [True,False] | True | True |


Metrics Comparison

Before and after hyperparameter tuning

Metrics Comparison

Non tuned performance

| | accuracy | roc_auc | f1 | recall |
|-----|----------|---------|-------|--------|
| SVC | 83.56 | 83.59 | 83.15 | 84.33 |
| lr | 81.37 | 81.41 | 80.95 | 82.34 |
| KNN | 78.90 | 79.01 | 78.85 | 81.77 |
| rf | 82.88 | 82.91 | 82.47 | 83.76 |

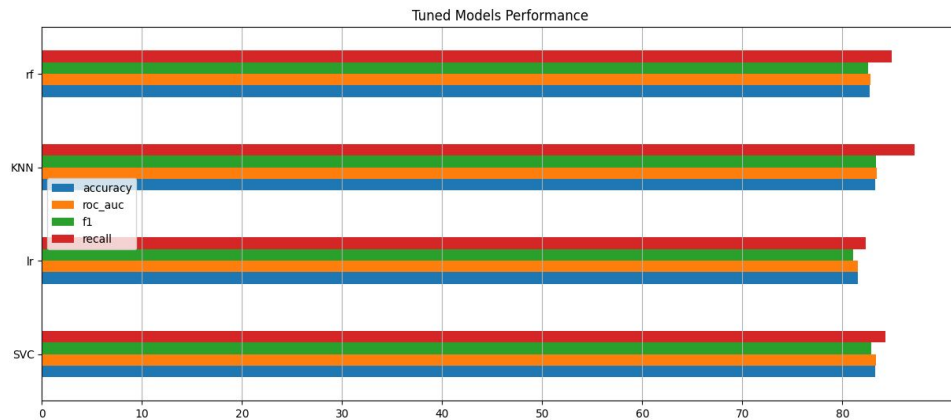


```
{'accuracy': 'rf', 'roc_auc': 'rf', 'f1': 'rf', 'recall': 'rf'}
```

Metrics Comparison

Tuned Performance

| | accuracy | roc_auc | f1 | recall |
|-----|----------|---------|-------|--------|
| SVC | 83.29 | 83.33 | 82.91 | 84.33 |
| lr | 81.51 | 81.54 | 81.07 | 82.34 |
| KNN | 83.29 | 83.43 | 83.38 | 87.18 |
| rf | 82.74 | 82.82 | 82.55 | 84.90 |



```
{'accuracy': 'SVC', 'roc_auc': 'KNN', 'f1': 'KNN', 'recall': 'KNN'}
```

UniGe

DIBRIS